

Website Knowledge & Context Readiness Audit

Twin Cities Gun & Pawn — OKF-inspired structure & agent-legibility grade

Industry: Firearms retail / pawnbroker (FFL dealer)

Compliance context: Firearms law (federal/state/local — disclaimer present), ADA/WCAG (relevant, not formally certified)

Audit type: Post-build re-score (after page build + schema v3.1 + canonical/sitemap/image-optimization work)

Note on OKF: Principles borrowed to judge a live site; this is **not** a claim of OKF compliance. Readiness ≠ ranking ≠ indexing.

Inspection scope

Assessed (read directly from disk — the exact files deployed via the build pipeline): all 18 production pages — `index`, `about`, `guns-rifles`, `accessories`, `pawn-loans`, `contact`, `faq`, `faq-gun-pawns`, `employment`, `resources`, `rules-for-pawning`, `gun-license-mn`, `unregistered-gun`, `gun-law-checklist`, `terms`, `privacy`, `sitemap`, `equal-opportunity` — plus `robots.txt`, `sitemap.xml`, and the generator sources (`build.py`, `build_pages.py`, `tools/schema_config.py`, `tools/build_schema.py`, `tools/check_schema.py`).

NOT assessed (state honestly):

- **Live production rendering** — the deployed host `twin-cities-pawn-guns.pages.dev` returned **HTTP 000** from the audit environment (sandbox IP cannot reach that Cloudflare host; `google.com` returned 200, confirming general connectivity). All scores below are from the **local built HTML**, which the build pipeline copies verbatim to deploy. Live-only signals — real HTTP response headers/`_headers` behavior, production cache TTLs, actual rendered CLS, server compression — are **not assessed**.
 - **Lighthouse field metrics** (Perf 86 / A11y 96 / Best-Practices 100 / SEO 100; CLS 0.181) are referenced **only** from the screenshots you supplied, not independently re-run here. They are cited as your evidence, not the audit's own measurement.
 - No prior **baseline** audit run was recorded, so a true per-category **delta cannot be computed**. Where build work clearly moved a category, it's noted qualitatively.
-

Overall grade: C+ — 75.5 / 100

Structurally excellent and highly machine-legible; held back by two concrete, fixable gaps — a **broken contact-form destination** and **zero agent/maintainer documentation** (no README / AGENTS.md / llms.txt) — plus one **entity-name inconsistency** between the visible metadata and the schema.

Score table

#	Category	Weight	Score	Element-specific reason
1	Metadata quality	15%	9/10	Every page has <pre><title> , <meta description> , self-referencing <link rel=canonical> , full Open Graph (og:title/de- scription/im- age/ type=business.b usiness/url/ locale/ site_name), Twitter card, and geo meta (geo.region=US- MN , geo.position=45 . 2619;-93.4499). JSON-LD v3.1 on all 18 pages with a 52-entry knowsAbout . -1: og:title / <title> say “Twin Cities Gun & Pawn” while schema name = “Twin Cities Pawn & Gun” — an agent reading both gets two entity names.</pre>
2	Content structure & machine-readability	20%	9/10	Clean semantic HTML: exactly one <h1> per page, ordered <h2> / <h3> , and real landmarks

#	Category	Weight	Score	Element-specific reason
				(<header> <nav> <main> <section>×10 <article>×4 <footer> on index). Meaning is in markup, not styling. An agent can parse without guessing.
3	Knowledge portability	15%	7/10	Fully static HTML — no proprietary/Javascript-locked surface; JSON-LD, sitemap.xml (18 URLs), robots.txt all present and exportable. -3: no llms.txt or small OKF-style services bundle; agent-facing knowledge lives only inside per-page JSON-LD, nothing portable at a single well-known path.
4	Cross-linking & navigation	15%	9/10	BreadcrumbList on all 18 pages; dense internal graph (index links to every page — contact ×8, guns-rifles ×7, pawn-loans ×6, resources cluster inter-linked); <nav aria-la-

#	Category	Weight	Score	Element-specific reason
				bel="Main navigation"> with dropdowns. Clear hierarchy Home → Inventory/Resources → leaf.
5	Human readability & UX	15%	8/10	99 of 102 <code></code> have descriptive <code>alt</code> ; the 3 empty-alt images are correctly decorative (<code>alt=""</code> <code>aria-hidden="true"</code> full-bleed CTA bg). All form controls use <code><label for></code> ; <code><meta viewport></code> present; images <code>loading="lazy"</code> . -2: CLS 0.181 (your screenshot) from web-font swap on the hero — a real layout-shift/UX cost; and <code>pawn-loans</code> meta description is 183 chars (truncates in SERP).
6	Documentation & maintainability	10%	3/10	Genuinely absent: no README , AGENTS.md / CLAUDE.md , or llms.txt anywhere in the repo (verified by find). A strong “edit generators,

#	Category	Weight	Score	Element-specific reason
				never generated HTML” convention exists in the code but is undocumented — a new coding agent would have to reverse-engineer it and risks editing generated files. <code>check_schema.py</code> is the one self-documenting guardrail.
7	Forms & interactive elements	10%	5/10	Contact form is well-built structurally — <code><label for> + id</code> on every field, <code>required</code>, correct <code>type=email/tel</code>, a <code><select></code> for subject. -5: <code>action="https://formspree.io/f/RE-PLACE_WITH_YOUR_ID"</code> is an unconfigured placeholder — the form submits nowhere; no data destination is documented, and there’s no visible client-side validation/confirmation

#	Category	Weight	Score	Element-specific reason
				beyond native required .

Weighted total: $(9 \times 1.5) + (9 \times 2.0) + (7 \times 1.5) + (9 \times 1.5) + (8 \times 1.5) + (3 \times 1.0) + (5 \times 1.0) = 75.5 / 100$

Top 3 strengths

1. **Schema-on-every-page, integrity-checked.** All 18 pages carry a JSON-LD `@graph` (`WebSite` → `PawnShop` → `WebPage` → `BreadcrumbList` + `Service/FAQ` nodes), and `tools/check_schema.py` enforces `@id` integrity and “Rule 11” traps on each build. That’s the single biggest agent-legibility asset — structured, validated, machine-queryable entity data.
2. **Textbook semantic + navigation structure.** One `<h1>`, ordered headings, real HTML5 landmarks, `BreadcrumbList` everywhere, and a dense internal link graph where every page is reachable and related pages cross-reference. An agent can build an accurate site relationship graph from the markup alone.
3. **Complete, consistent metadata surface.** Canonical + full Open Graph + Twitter + geo meta on every page, with matching OG image dimensions (1200×630) and alt text — clean social/agent preview data with no missing core fields.

Top 5 prioritized fixes

1. **Wire up (or remove) the contact form destination.** Problem: `contact.html` action is `forms-pree.io/f/REPLACE_WITH_YOUR_ID` — submissions go nowhere. Why it matters (humans + agents): the primary conversion path is silently broken, and an agent inspecting the form cannot determine a real data destination. Action: replace with the real Formspree (or other) endpoint in the generator, or swap to a documented `mailto: /backend`; then note the destination in a form comment.
2. **Reconcile the business name across metadata and schema.** Problem: visible title/OG = “Twin Cities Gun & Pawn”; schema `name` = “Twin Cities Pawn & Gun”. Why it matters (agents): two names for one entity weakens entity resolution and citation confidence. Action: pick one canonical name (with the other as schema `alternateName`) and make `schema_config.py` + page titles agree.
3. **Add an `llms.txt` (and a tiny services bundle).** Problem: no single well-known machine-readable summary of who/what/where. Why it matters (agents): `llms.txt` is the emerging front door for AI agents to cite services, hours, address, and page map without scraping 18 pages. Action: generate `/llms.txt` from `schema_config.py` (name, address, phone, hours, services, key page URLs) as a build step.
4. **Add `README.md` + `AGENTS.md` documenting the build pipeline.** Problem: the “edit generators, not generated HTML; run `build_pages` → `build_schema` → `check_schema`” workflow is undocumented. Why it matters (agents + maintainers): a coding agent will otherwise edit generated `.html` and have changes silently overwritten on rebuild. Action: a short `AGENTS.md` stating the golden rule, the build commands, and the source-of-truth files.

5. **Cut CLS and tidy two metadata edges.** Problem: CLS 0.181 (hero font swap) exceeds the 0.1 “good” threshold; `pawn-loans` description is 183 chars; homepage canonical targets `/index.html` rather than `/`. Why it matters (humans + ranking-adjacent): layout shift hurts UX and Core Web Vitals; over-length descriptions truncate. Action: preload the hero font / reserve space with `font-display:optional` or sized containers; trim the description to ≤ 160 chars; canonicalize the homepage to the bare domain.

Verdict

The site is genuinely agent-ready at the content layer — validated JSON-LD on every page, clean semantic structure, breadcrumbs, and a dense internal link graph mean an AI agent can parse the entity, its services, and its geography without guessing. What keeps it from an A is not structure but two operational gaps and one consistency slip: a **contact form that submits to an unconfigured placeholder**, **no agent/maintainer documentation** (README/AGENTS.md/llms.txt), and a **name mismatch between the visible metadata and the schema**. The single highest-impact next change is **fixing the contact-form destination** — it’s the site’s main conversion path and currently silently fails — immediately followed by **adding an `llms.txt` + `AGENTS.md`**, which together would lift Documentation and Portability and push the overall grade into the mid-B range on the next re-score.

Re-run this audit against the live URL once the sandbox (or your machine) can reach `twin-cities-pawn-guns.pages.dev` to capture the live-only signals marked “not assessed.”