

Indexing Diagnostic & Fix Runbook — Twin Cities Gun & Pawn

Site: <https://twin-cities-pawn-guns.pages.dev>

Repo: MMG001/Twin-Cities-Pawn-Guns (branch `main`)

Date: September 30, 2026

Symptom addressed: pages sitting in Google Search Console as “**Crawled — currently not indexed.**”

TL;DR — the honest verdict

I ran your runbook as a **measured diagnostic** against the actual built HTML, step by step. The headline result:

This is NOT a duplicate-content problem. The runbook is written mainly for sites with dozens of near-identical city / location landing pages. Your site is a **single-location business with topically distinct pages**, and the measurements confirm the pages are genuinely different once boilerplate is removed. **So the fix is NOT to rewrite your pages** — doing that would risk the content that is already good.

What the measurements did surface as legitimately worth fixing was **thin contextual internal linking** on the resource guides (Step 2). That has been fixed and pushed this session. The technical indexability checks (Step 3) all came back **clean**.

Two honesty caveats up front, because they bound everything below:

1. **I have no access to your Google Search Console.** I cannot see which specific URLs Google flagged, cannot read the Page Indexing report, and cannot submit URLs for indexing. Those steps (Step 4) are things **you** must do in GSC — I’ve written them out precisely.
 2. **The live site was unreachable from this sandbox** (network returns HTTP 000). So every check below was run against the **built HTML in the repo** — which is the exact source that Cloudflare Pages deploys, but it is not a live fetch of the deployed page. Where that distinction matters, I say so.
-

Step 0 — Measure duplication (is duplication even the cause?)

The runbook’s first instruction is to measure before assuming. I did two passes.

Pass 1 (naive): full visible text of each page, 5-word shingles, Jaccard similarity. This produced high numbers — but they were **inflated by boilerplate** (the shared nav, ticker, footer, and hours/location block appear on every page). That is a measurement artifact, not real duplication.

Pass 2 (correct): isolate only the `<main>` content body, strip city/geo tokens, then compute 5-word-shingle Jaccard similarity between every pair of content pages. This is the number that actually matters.

Result — pairwise content similarity (main body only):

Page	Similarity vs. rest of site
about.html	~8.2%
rules-for-pawning.html	~8.4%
gun-license-mn.html	~6.9%
unregistered-gun.html	~5.5%
gun-law-checklist.html	~6.5%
faq-gun-pawns.html	~10.5%
pawn-loans.html	~19.1%
index.html	~22.9%

Runbook threshold: below ~40% overlap → duplication is not the cause. **Every page is far below that.** The two highest (index and pawn-loans) are elevated only because they legitimately restate the core service vocabulary (“pawn loan”, “firearm”, “Twin Cities”) — which is expected and healthy, not duplication.

Thin pages noted (by design, not a defect): sitemap (164 words), equal-opportunity (187), employment (211), resources (253), privacy (265). These are utility/legal pages; being short is normal for their type and is not what triggers “crawled — not indexed” on content pages.

Conclusion for Step 0: Skip Step 1 (make pages unique) entirely — the pages are already unique. Proceed to Steps 2 and 3.

Step 1 — Make pages unique

Skipped, deliberately. Step 0 proved the pages are already distinct. Rewriting them would be effort spent against a non-problem and would put good content at risk. No changes made here.

Step 2 — Strengthen internal links from indexed pages



This is the one place the measurements found a real, fixable weakness.

How I measured it: for every page, count inbound links, separating (a) boilerplate inbound (nav + footer, which every page gets automatically) from (b) contextual inbound — links that sit inside another page’s `<main>` content. Contextual links are the ones Google weighs most for discovering and valuing a page.

Finding: no orphan pages (every page had 17 boilerplate inbound links), **but** several resource guides had very thin contextual inbound linking — Google was reaching them only through the nav/footer, which is a weak signal.

Fix applied: added a contextual “**Related Pages**” module to the resource guides, generated from the source (`build_pages.py` → `info_page()`), so it survives regeneration. Mappings:

- **rules-for-pawning** → Pawn Loans, Gun License in MN, Gun Pawn FAQ
- **gun-license-mn** → Guns & Rifles, Rules for Pawning, Unregistered Firearms
- **unregistered-gun** → Gun License in MN, 2026 MN Gun Law Checklist, Guns & Rifles

Measured before → after (contextual inbound links):

Page	Before	After
guns-rifles.html	5	7
gun-license-mn.html	4	6
pawn-loans.html	5	6
rules-for-pawning.html	5	6
unregistered-gun.html	2	3
faq-gun-pawns.html	2	3
gun-law-checklist.html	—	4

Verification: regenerated all 18 pages, `check_schema.py` → **ALL PASS (18/18)**; confirmed the “Related pages” section renders on all three guides; re-ran the link-count script to produce the “after” column above. Committed and pushed (`0e79a22`).

Step 3 — Technical indexability check ALL CLEAN

Every technical blocker the runbook lists was checked against the built HTML:

Check	Result
<code>noindex</code> meta tag anywhere	None — no page contains <code>noindex</code>
X-Robots-Tag: <code>noindex</code> in <code>_headers</code>	None — headers file is clean
robots.txt disallow rules	None blocking — robots.txt allows all crawlers; points to sitemap
<code><meta name="robots"></code> per page	Present = <code>index, follow</code> on all 19 built pages
Self-referencing canonical	Present & correct on every page
Content in raw HTML (not JS-injected)	Confirmed — full text is in server HTML (e.g. <code>index <main></code> ≈ 5,060 chars). JavaScript only defers the Google Maps iframe (<code>main.js</code>) — no content depends on JS
sitemap.xml accuracy	18 URLs = 18 production pages ; <code>font-mockup.html</code> correctly excluded

Conclusion: there is no technical reason Google can't index these pages. "Crawled — currently not indexed" here is a **quality/priority judgment by Google**, not a crawl block — which is consistent with the Step 0 finding (unique but modest-authority pages on a young site).

One structural nuance for you to decide (not forced)

The homepage is reachable at **both** `https://.../` and `https://.../index.html` :

- Schema `WebSite.url` uses the bare `/`
- The homepage canonical, schema `@id`, and sitemap entry use `/index.html`
- 58 internal links point to `href="index.html"`

Google can crawl the home content at two URLs. This is **not** currently causing an indexing failure, but canonicalizing the homepage to one form (bare `/` is the conventional choice) is a tidy long-term improvement. Because it's a structural URL change touching ~58 links, the sitemap, and schema `@id` s, **I did not make it unilaterally** — tell me if you want it and I'll do it cleanly in one pass.

Step 4 — Submit for indexing (YOUR steps — I cannot do these)

I have **no GSC access**, so these are for you to run in Google Search Console for the `twin-cities-pawn-guns.pages.dev` property:

1. **Confirm the flagged URLs.** Open **Indexing** → **Pages** → "Crawled — currently not indexed" and note exactly which URLs are listed. (I've been diagnosing the whole site; this tells us if it's the resource guides specifically.)
2. **Re-submit the sitemap.** Indexing → Sitemaps → submit `https://twin-cities-pawn-guns.pages.dev/sitemap.xml` (it now reflects the current 18 pages).

3. **URL Inspection + Request Indexing** for each flagged URL, after the new deploy is live (verify the “Related Pages” section is visible on the deployed guide first). Do a handful per day; don’t mass-spam requests.
4. **Prioritize** the pages that now have stronger internal links (gun-license-mn, unregistered-gun, rules-for-pawning) — the improved contextual linking is a fresh signal for Google to re-evaluate them.

Step 5 — Monitor (YOUR ongoing steps)

- Give Google **2-4 weeks**; “Crawled — not indexed” reversals happen on Google’s crawl schedule, not on demand.
- Watch the GSC Pages report for URLs moving from “Crawled — not indexed” → “Indexed.”
- If specific pages are still stuck after a month **and** GSC shows them under a “Duplicate” bucket (not the generic “crawled” one), come back and we’ll do a targeted content-differentiation pass on just those URLs — but only with GSC evidence pointing there.

What was changed this session (for your records)

Commit	What
5c2f20e	Canonical name → “Twin Cities Gun & Pawn” across schema, titles, OG, page copy (prior turn)
0e79a22	Contextual “Related Pages” internal-linking module added to resource guides (this runbook, Step 2)

All edits were made to the source generators (`build_pages.py`), then regenerated — so they survive future rebuilds. No generated `.html` was hand-edited.

Honest limits of this report

- **No Google Search Console access** → I cannot confirm which URLs Google flagged, cannot read the Page Indexing report, and **cannot submit or “request indexing.”** Steps 4-5 are yours.
- **Live site unreachable from this environment** (HTTP 000) → all checks ran against the **built HTML in the repo**, which is the exact deploy source, but is not a live fetch of the deployed page. Re-verify the deployed pages in a browser after Cloudflare finishes building commit `0e79a22` .
- **I am not claiming the indexing issue is “fixed.”** Reversing “Crawled — currently not indexed” is Google’s decision and takes crawl cycles. What I can verifiably claim: duplication was measured and ruled out, internal linking was measurably strengthened, and no technical indexability blocker exists in the source.